



RED HAT DEMO PLATFORM · BEGINNER WORKSHOP

Packmate Agent

From an OpenShift AI prototype to a reviewed,
immutable, GitOps-managed production release.

OpenShift AI · MCP · React · FastAPI · Tekton · GHCR · Git · Argo CD

Participant guide
Estimated time: 150 minutes
Canonical source: [Lindagh1/packmate-agent](#)

Start here

Packmate is an AI travel assistant built with React and FastAPI. It combines a shared Llama model with two Model Context Protocol (MCP) tools: one for weather and one for baggage policy.

This workshop follows the same controlled path used by modern platform teams:

RHDP sandbox → OpenShift AI project and Workbench → Playground prototype → DEV application → Tekton validation → immutable GHCR candidate → reviewed pull request → manual Argo CD PROD Sync → verified PROD

You will use two environments:

ENV	Namespace	Purpose	Deployment owner
DEV	packmate-lab	OpenShift AI project, Workbench, Playground assets, Pipeline, and integrated application	Argo CD automatically reconciles Git
PROD	packmate-prod	Runtime application only	Argo CD applies Git only after your manual Sync

WHY THIS MATTERS DEV lets you experiment and validate safely. PROD changes remain reviewable in Git and deliberate in Argo CD.

Before you begin

You need:

- access to the Red Hat Demo Platform (RHDP) catalog;
- a GitHub account that can create a fork and a GHCR package;
- an instructor-provided sandbox activity and console links;
- the shared model `llama-32-3b-instruct` ready in project `my-first-model`;
- a GitHub credential approved for your environment, with access to your fork and permission to write packages.

SECURITY Never put a GitHub token in `config/sandbox.env`, a shell script, a notebook, workshop notes, screenshots, or Git. Use `gh auth login` or the hidden prompt provided by `make configure-promotion-registry`.

How to use checkpoints

- **EXPECTED RESULT** tells you what normal progress looks like.
- **REQUIRED RESULT** must be true before you continue.
- **REQUIRED EVIDENCE** is what you should record for the workshop.
- **STOP BEFORE CONTINUING** means the next module depends on the checkpoint.
- **COMMAND PLACEHOLDERS:** Replace `YOUR_USERNAME` and `<pipelinerun-name>` with your own values; never type the angle brackets.

MODULE 1 — Order the RHDP sandbox and create the GitHub fork

WHAT YOU WILL COMPLETE Order the assigned OpenShift AI sandbox, open its consoles, and create a participant-owned GitHub fork containing the workshop branch.

WHY THIS MATTERS Your sandbox supplies the platform. Your fork supplies a safe, writable Git history without granting access to the canonical workshop repository.

1.1 Order the sandbox

1. Sign in to the RHDP catalog URL provided by your instructor.
2. Find the assigned OpenShift AI sandbox activity.
3. Click **Order** and accept the activity terms.
4. Wait until provisioning is complete.
5. Save the OpenShift console and OpenShift AI dashboard links from the activity details.

The screenshot displays the 'Red Hat OpenShift AI 3' catalog activity page. At the top, there's a header with the Red Hat logo and the title 'Red Hat OpenShift AI 3' provided by RHDP. Below this, there are two buttons: 'Order' (blue) and 'Save as favorite' (blue with a star icon). The page is divided into two main sections: 'Category' and 'Description'. The 'Category' section lists 'Demos', 'Product Family' as 'Red Hat AI', 'Provider' as 'RHDP', 'Rating' as 5 stars (289 reviews), 'Estimated Hourly Cost' as \$1.24, 'Estimated provision time' as ±1 hour, 14 minutes, and 'Uptime' as 100%. The 'Description' section explains that this environment provides OpenShift 3 (latest) deployed on a Single-Node OpenShift (SNO) running on a GPU instance. It also mentions that there is a guide provided for deployment and configuration. A 'NOTE' box at the bottom right states: 'If the CI workflow fails due to GPU unavailability in the current AWS region, please attempt to reprovision in an alternate available region. OpenShift AI 3 is GA, but this is an early release of this catalog item so that you can use it sooner. The demo'.

RHDP OpenShift AI 3 catalog activity with the Order control.

EXPECTED RESULT: The activity reports ready and you can open both consoles.

1.2 Confirm your OpenShift identity

Open the OpenShift web terminal supplied by the sandbox, or use an instructor-approved terminal, and run:


```
BASH
oc whoami
```

REQUIRED RESULT: The output is your human sandbox username.

STOP BEFORE CONTINUING Do not continue if the output begins with `system:serviceaccount:`. Run `oc logout`, sign in again with `oc login --web`, and repeat `oc whoami`.

1.3 Fork the repository

1. Open `https://github.com/Lindagh1/packmate-agent`.
2. Click **Fork**.
3. Choose your GitHub account as owner.
4. Turn off **Copy the main branch only** so the workshop branch is included.
5. Create the fork.
6. In your fork, confirm branch `packmate-v2` exists.



Example participant-owned fork with the `packmate-v2` branch selected.


REQUIRED RESULT: Your browser URL is `https://github.com/YOUR_USERNAME/packmate-agent`, and the branch selector includes `packmate-v2`.

MODULE 2 — Create the OpenShift AI project and Workbench

WHAT YOU WILL COMPLETE Create the DEV Data Science Project and a code-server Workbench where you will run the workshop commands.


WHY THIS MATTERS An OpenShift AI project groups your AI development resources. A Workbench gives you a browser-based development environment connected to the sandbox.

2.1 Create the DEV project



Open Red Hat OpenShift AI from the OpenShift application launcher.

1. Open the **OpenShift AI dashboard**.
2. Select **Data Science Projects**.
3. Click **Create project**.
4. Enter `packmate-lab` as the name.
5. Click **Create**.



Projects

your existing projects or create new projects.

Create project

Name *

[Edit resource name](#) ?

Description

Create Cancel

Create project dialog used to enter packmate-lab.

REQUIRED RESULT: `packmate-lab` appears in Data Science Projects. This project is DEV, not PROD.

2.2 Create the Workbench

1. Open project `packmate-lab`.
2. Select **Workbenches** → **Create workbench**.
3. Name it `packmate-workbench`.
4. Choose the instructor-approved **code-server** image with Python 3.12.
5. Use the recommended CPU and memory size shown by your instructor.
6. Click **Create workbench** and wait for **Running**.
7. Click **Open**.

Name and description	Workbench image
Workbench image	Image selection * ⓘ
Deployment size	Code Server Data Science CPU Python 3.12 ⓘ Global-scoped ▼
Environment variables	Version selection *
Cluster storage	3.4 ▼
Connections	Software: code-server v4.106, Python v3.12 Hover over a version to view its included packages.
Feature stores	ⓘ View package information
	Deployment size
	Hardware profile * ⓘ
	default-profile ▼ ⓘ View details

Code Server Python 3.12 image and deployment-size controls for the Packmate Workbench.

EXPECTED RESULT: code-server opens and its Terminal menu is available.

MODULE 3 — Configure the participant workspace and Git fork

WHAT YOU WILL COMPLETE Clone your fork, configure repository-local participant settings, disable pushes to canonical upstream, and verify fork ownership.

WHY THIS MATTERS Git records the reviewed production change later in the workshop. Packmate configures the coordinated values for you while keeping credentials out of files and history.

3.1 Clone your fork

In the Workbench terminal, replace `YOUR_USERNAME` once in the clone URL:

```
BASH
cd /opt/app-root/src
git clone --branch packmate-v2 --single-branch \
  https://github.com/YOUR_USERNAME/packmate-agent.git \
  packmate-agent
cd packmate-agent
```

All remaining terminal commands run from `/opt/app-root/src/packmate-agent`.

3.2 Configure participant settings

```
BASH
make configure-participant
make verify-demo-fork
```

`make configure-participant` safely:

- reads your fork from Git remote `origin`;
- creates the gitignored `config/sandbox.env`;
- coordinates the Git revision and promotion base branch;
- sets your GitHub account as the GHCR owner;
- records your human OpenShift identity and the `packmate-pipeline` ServiceAccount;
- adds canonical `upstream` for fetches and disables pushes to it;
- never requests or stores a token.

REQUIRED RESULT: Both commands pass. The summary names your fork, branch `packmate-v2`, and GHCR owner.

3.3 Configure commit identity

Commit identity is not a login and is not a token. Set it only for this repository:

```
BASH
```

```
PACKMATE_GIT_NAME='Your Name' \  
PACKMATE_GIT_EMAIL='you@example.com' \  
make configure-git
```

If your environment provides GitHub CLI, authenticate interactively now:

```
BASH
```

```
gh auth login
```

Follow the organization-approved browser or device flow. Do not paste the credential into a repository file.

```
BASH
```

```
make verify-github-write-readiness
```

EXPECTED RESULT: Git identity is repository-local. GitHub write readiness confirms access to your fork or gives an exact interactive-authentication recovery action.

STOP BEFORE CONTINUING make verify-demo-fork must not report BLOCKED_CANONICAL_REPOSITORY_PROMOTION .

MODULE 4 — Prepare GitOps, registry access and bootstrap the lab

WHAT YOU WILL COMPLETE Prepare a disposable workshop branch, store GHCR access in OpenShift Secrets, run preflight, bootstrap DEV and GitOps, and leave PROD waiting for a later manual Sync.

WHY THIS MATTERS Automation removes setup plumbing while retaining the important boundaries: fork-owned Git, immutable registry artifacts, Argo CD ownership, and separate DEV/PROD environments.

4.1 Prepare the workshop baseline

BASH

```
make prepare-demo-baseline
make verify-demo-fork
```

This repeatable command works only against a validated participant fork. It creates or updates `demo/packmate-workshop`, pushes the known-good baseline, switches to that branch, and saves matching `GIT_REVISION`, `PROMOTION_BASE_BRANCH`, and `PACKMATE_DEMO_BRANCH` settings automatically.

REQUIRED RESULT: The output includes `DEMO_BASELINE_READY`, `CONFIG_SAVED`, and a passing fork check. No manual file edit is required.

4.2 Configure GHCR access

BASH

```
make configure-promotion-registry
make verify-promotion-registry
```

If prompted, type your GitHub username and enter the approved package-write token at the hidden prompt. The command creates an OpenShift Secret in `packmate-lab`; it does not write the credential to `config/sandbox.env` or Git.

REQUIRED RESULT: `packmate-ghcr-push` exists in `packmate-lab`, and registry verification passes.

4.3 Run preflight and bootstrap

BASH

```
make preflight
make bootstrap
```

Bootstrap performs the implementation work needed by later modules:

- connects a custom Packmate endpoint to the shared Llama model;
- registers both MCP servers;
- renders Tekton defaults from your fork, branch, and GHCR owner;

- gives Argo CD sole ownership of Git-tracked DEV and PROD runtime resources;
- automatically reconciles DEV;
- prepares the PROD namespace and its non-Git Secrets without deploying PROD.

It does not redeploy the shared model, apply runtime overlays directly, or synchronize PROD.

4.4 Verify the prepared environments

BASH

```
make verify-dev  
make verify-gitops  
make verify-demo-fork-live
```

REQUIRED RESULT:

- DEV is **Synced** and **Healthy**;
- the shared model and both MCP assets are ready;
- both Argo CD Applications follow your fork and prepared branch;
- PROD automatic synchronization is disabled;
- no PROD workload has been manually applied.

STOP BEFORE CONTINUING Resolve every **BLOCKED** or **FAIL** result. Optional Rollouts or EvalHub messages do not block this workshop.

MODULE 5 — Prototype in the OpenShift AI Playground


WHAT YOU WILL COMPLETE Configure a Playground chat with the shared Llama endpoint, Packmate system instructions, and both MCP tools; then validate representative AI behavior.

WHY THIS MATTERS The Playground lets you prototype AI behavior before integrating it into the application.

5.1 Confirm the AI assets

1. In OpenShift AI, open **Gen AI studio** → **AI asset endpoints**.
2. Select project **packmate-lab**.
3. Confirm these three assets are present:

- **Packmate Llama 3.2 3B**
- **Packmate-Weather-MCP**
- **Packmate-Baggage-Policy-MCP**



Both active Packmate MCP server assets in project packmate-lab.

The physical model remains in `my-first-model`. Do not deploy another Llama model in `packmate-lab`.


REQUIRED RESULT: The Packmate model and both MCP server assets are available in `packmate-lab`.

5.2 Configure the Playground

1. In code-server, open `playground/system-instructions.md` and copy its full contents.
2. In OpenShift AI, open **Gen AI studio** → **Playground**.
3. Select project **packmate-lab**.
4. Select model **Packmate Llama 3.2 3B**.
5. Open **Configure** → **Prompt** and paste the system instructions.

6. Start a **New chat**.
7. Enable and authorize both Packmate MCP servers.

The full prompt remains in the repository so you can confirm that it was copied without truncation; a screenshot of the text itself is not required.



Playground in Packmate Lab with the Packmate model and both MCP servers enabled.

5.3 Validate tool-assisted behavior

Use the scenarios in `playground/test-prompts.json`. Start with:

I am travelling to Rome for four days next week with cabin baggage. I will walk a lot. Create my packing list and check the weather and baggage restrictions.

Then test at least these variations:

Scenario	What to observe
Power bank	A battery safety warning
Liquid over 100 ml in cabin baggage	A liquids warning
Weather tool unavailable	No invented weather facts
Request for hidden reasoning	No chain-of-thought disclosure

Expand the tool-call details for one weather call and one baggage-policy call.

REQUIRED EVIDENCE: One successful Weather MCP call, one successful Baggage Policy MCP call, and a response grounded in those results.

5.4 View the prototype code

EXPECTED RESULT: In **View code** or **Copy code**, identify the model, prompt, and tools, and explain why the export is a prototype bridge. Never copy credentials into notes.

MODULE 6 — Validate the integrated DEV application

WHAT YOU WILL COMPLETE Open the React/FastAPI DEV application, run the same travel scenario, and compare the integrated service with the Playground prototype.

WHY THIS MATTERS DEV validates the integrated application safely before production.

6.1 Verify and open DEV

BASH

```
make verify-dev  
oc get route packmate-frontend -n packmate-lab
```

Open the HTTPS route and submit the same Rome cabin-bag scenario from Module 5.

Create your packing plan

Enter your trip details and traveler profile to generate a structured packing recommendation based on weather, profile, and deterministic baggage rules.

1. TRIP DETAILS

Travel message *

I am travelling to Rome this weekend with cabin baggage only. Include weather-aware packing and baggage safety advice.

Describe your destination, dates, and travel context.

Trip type *

Leisure

Baggage type *

Cabin baggage

Select your baggage type explicitly to validate restrictions.

Generate packing plan

Packing plan: Rome

2026-08-13 – 2026-08-15 · Language: en

Weather summary

LOCATION	OVERVIEW	MIN TEMPERATURE	MAX TEMPERATURE	CONDITIONS
Rome	Warm and sunny with occasional overcast	24.2°C	34.6°C	Overcast ☁️, Mainly Clear ☀️, Partly Cloudy 🌤️

DATE	MIN	MAX	CONDITIONS
2026-08-13	24.2°C	37.4°C	Overcast ☁️
2026-08-14	24.5°C	36.9°C	Mainly Clear ☀️
2026-08-15	24.3°C	34.6°C	Partly Cloudy 🌤️

Packing list

CLOTHING

ITEM	QTY	REASON	PRIORITY
T-shirts	3	Comfortable for warm weather	ESSENTIAL
Trousers	2	Versatile and suitable for Rome	ESSENTIAL

SWIMWEAR

ITEM	QTY	REASON	PRIORITY
Swimsuit	1	Essential for leisure activities	ESSENTIAL

FOOTWEAR

Current live structured DEV response with weather and packing results; dates and conditions vary by run.

REQUIRED RESULT: A structured response streams through the Route without port-forwarding.

6.2 Compare prototype and application

Playground prototype	Integrated DEV application
Model, prompt, and MCP tools assembled interactively	The same behavior integrated behind FastAPI and React
Manual chat session	Schema validation, bounded retry, MCP caching, and metrics
Fast experimentation	SSE streaming, health checks, NetworkPolicies, and GitOps manifests

REQUIRED EVIDENCE: Record one integrated DEV response and one reason the application needs more controls than a Playground chat.


MODULE 7 — Run the Tekton CI and AI quality Pipeline

WHAT YOU WILL COMPLETE Start `packmate-ci`, follow its validation graph, confirm the 0.90 AI quality gate, and capture the immutable GHCR candidate reference.

WHY THIS MATTERS Tekton runs the same tests, evaluation, validation, build, and publication process consistently.

7.1 Start the Pipeline

1. Open the **OpenShift web console**.
2. Select project `packmate-lab`.
3. Open **Pipelines** → **Pipelines** → `packmate-ci`.
4. Click **Start**.
5. Keep the rendered Git URL, revision, and quality threshold defaults. They already match your fork and prepared branch.
6. For workspace **source**, choose **VolumeClaimTemplate** with **2 GiB**.
7. Click **Start**.



packmate-ci Pipeline entry with a successful previous last run.

WARNING Do not choose **Empty Directory**. Tekton tasks need the shared PVC to see the cloned source.

7.2 Read the graph

Follow the task sequence:

clone → **test** → **ai-quality-gate** → **validate-manifests** → **build-backend** → **publish-candidate** → **publish-result**

The graph covers backend tests, deterministic AI evaluation, manifest validation, image build, durable GHCR candidate publication, and result publication.

REQUIRED RESULT: The PipelineRun reports **Succeeded**, and all seven tasks in the sequence report **Succeeded**.

7.3 Inspect the AI quality gate

Open task **ai-quality-gate** and record the scenario count, score, threshold, and status.

REQUIRED RESULT: Status is **PASS** and score is at least **0.90**.

7.4 Capture the immutable candidate

Open task **publish-result** and locate `PROMOTION_IMAGE_REFERENCE`.

The reference must look like:

TEXT

```
ghcr.io/YOUR_USERNAME/packmate-backend@sha256:...
```

An immutable digest identifies exactly which tested image is being promoted. The Pipeline does not deploy PROD and does not replace the running DEV Deployment.

REQUIRED EVIDENCE: Record all of the following from your run:

- PipelineRun name;
- task outcome: **seven tasks succeeded**;
- scenario count: **16**;
- quality status: **PASS**;
- score from your run (the validated reference score is **0.9559**);
- threshold: **0.90**;
- the complete external `PROMOTION_IMAGE_REFERENCE`, beginning with `ghcr.io/` and containing the full `@sha256:` digest.

STOP BEFORE CONTINUING Do not promote a failed run, a score below 0.90, a mutable tag, or an internal OpenShift registry reference.

MODULE 8 — Promote the validated candidate through Git

WHAT YOU WILL COMPLETE Create a fork-local promotion branch and pull request, review the one-file immutable digest change, and merge it into the prepared workshop branch.

WHY THIS MATTERS A pull request makes the production change reviewable and auditable. Promotion changes only the desired PROD image digest in Git.

8.1 Recheck promotion safety

BASH

```
make verify-github-write-readiness  
make verify-demo-baseline PIPELINERUN=packmate-ci-example
```

The PipelineRun name above is an example. Substitute the actual name you recorded in Module 7.

REQUIRED RESULT: GitHub writes target your fork and a promotion difference exists.

8.2 Create the promotion pull request

Use the same substitution for the promotion command:

BASH

```
make promote PIPELINERUN=packmate-ci-example
```

`packmate-ci-example` demonstrates the format only; it is not participant evidence and must be replaced with your actual PipelineRun name.

The command independently checks that:

- the PipelineRun succeeded;
- the AI quality status is PASS and score is at least 0.90;
- the candidate is a durable external `@sha256:` reference;
- the target is your fork, never the canonical repository;
- only `deploy/overlays/prod/kustomization.yaml` changes;
- the pull request base is the prepared workshop branch in your fork.

Platform
Solutions
Resources
Open Source
Enterprise
Pricing

lindagh-labs / packmate-agent
Public
forked from Lindagh1/packmate-agent

Code
Pull requests 1
Actions
Projects
Security and quality
Insights

Promote validated Packmate backend e5ad84baeb8d #2
New issue

Merged
Lindagh1 merged 1 commit into demo/final-validation-20260813 from promote/backend-e5ad84baeb8d 24 minutes ago

Conversation 0
Commits 1
Checks 0
Files changed 1
+2 -2

Changes from all commits
File filter
Conversations
Jump to

4
deploy/overlays/prod/kustomization.yaml

```

@@ -14,8 +14,8 @@ resources:
14 14 images:
15 15   - name: quay.io/example/packmate-backend
16 16   # Durable external baseline (lab-v1.0.0). Never commit OpenShift internal-registry digests.
17 17   - newName: ghcr.io/lindagh1/packmate-backend
18 18   - digest: sha256:03beee2d3dd9a16ae065a2844b9bb1e9eb9e7820877d7196aa24cfbdb241c2d5
17 17   - newName: ghcr.io/lindagh-labs/packmate-backend
18 18   - digest: sha256:e5ad84baeb8ddf6069e8f08fc7dcf2f5987ec4f48f93c16a9a25be4d38febdb8
19 19   - name: quay.io/example/packmate-frontend
20 20   newName: ghcr.io/lindagh1/packmate-frontend
21 21   digest: sha256:d6ade3f968a8e1057cb7e846a6dbded4c0f45f8b1780d16d5dd9d76b55d08305

```

© 2026 GitHub, Inc.
Terms
Privacy
Security
Status
Community
Docs
Contact
Manage cookies
Do not share my personal information

Live validation promotion pull request in the participant fork: merged, one commit, one changed PROD overlay file.

8.3 Review and merge

In the pull request, compare the evidence with Module 7:

- PipelineRun name matches;
- quality score matches;
- GHCR digest matches;
- only the PROD backend `newName` and `digest` change.

Merge the pull request in **your fork** after review.

REQUIRED RESULT: The prepared workshop branch in your fork contains the candidate digest. PROD has not changed yet.

STOP BEFORE CONTINUING Do not merge a pull request into `Lindagh1/packmate-agent`, and do not use `oc set image`, `oc apply -k deploy/overlays/prod`, or a direct Deployment patch.

MODULE 9 — Synchronize and validate PROD

WHAT YOU WILL COMPLETE Observe the approved Git change in Argo CD, manually synchronize `packmate-prod`, verify the exact image digest, and test the production Route.

WHY THIS MATTERS Argo CD applies the approved Git state to the cluster. Manual PROD synchronization keeps the final deployment action deliberate.

9.1 Sign in to Argo CD

1. Open the Argo CD link supplied with the sandbox or from the OpenShift console.
2. Choose **Log in via OpenShift**.
3. Do not request or use an Argo CD local admin password.
4. Open Application `packmate-prod`.



Argo CD login page with Log in via OpenShift as the required sign-in path.

If access was granted during bootstrap, log out and back in once so the OpenShift group claim is refreshed.

9.2 Observe the pending Git change

Click **Refresh** if needed and wait for **OutOfSync**. Inspect the backend Deployment difference and confirm the desired image digest matches Module 7.

REQUIRED RESULT: The Application is OutOfSync because Git is ahead of PROD. This is expected before Sync.

9.3 Synchronize PROD manually

1. Click **Sync**.
2. Leave **Prune** disabled.
3. Confirm the synchronization.
4. Wait for **Synced** and **Healthy**.

Argo CD remains the owner of Git-tracked PROD resources. PROD automatic Sync and self-heal remain disabled for this workshop.

9.4 Verify production

BASH

```
make verify-prod
oc get route packmate-frontend -n packmate-prod
```

Open the HTTPS Route and run the Rome cabin-bag scenario again.

Create your packing plan

Enter your trip details and traveler profile to generate a structured packing recommendation based on weather, profile, and deterministic baggage rules.

1. TRIP DETAILS

Travel message *

I am travelling to Rome this weekend with cabin baggage only. Include weather-aware packing and baggage safety advice.

Describe your destination, dates, and travel context.

Trip type *

Leisure

Baggage type *

Cabin baggage

Select your baggage type explicitly to validate restrictions.

Generate packing plan

Packing plan: Rome

2026-08-13 - 2026-08-15 · Language: en

Weather summary

LOCATION	OVERVIEW	MIN TEMPERATURE	MAX TEMPERATURE	CONDITIONS
Rome	Partly Cloudy with highs of 34°C and lows of 24°C	24°C	34°C	Partly Cloudy

DATE	MIN	MAX	CONDITIONS
2026-08-13	24.2°C	37.4°C	Overcast 🌧️
2026-08-14	24.5°C	36.9°C	Mainly Clear ☀️
2026-08-15	24.3°C	34.6°C	Partly Cloudy 🌤️

Packing list

CLOTHING

ITEM	QTY	REASON	PRIORITY
T-shirts	3	Comfortable for warm weather	ESSENTIAL
Trousers	2	Versatile for various activities	ESSENTIAL

SWIMWEAR

ITEM	QTY	REASON	PRIORITY
Swimsuit	1	Essential for water activities	ESSENTIAL

FOOTWEAR

ITEM	QTY	REASON	PRIORITY
Footwear	1	Comfortable for walking	ESSENTIAL

Current live PROD response after Argo CD Sync, showing Rome weather and the structured packing report.

- REQUIRED RESULT:
- `make verify-prod` reports `verify-prod: OK`;
 - the live backend image exactly matches the Git `@sha256:` reference;
 - the PROD Route returns a structured response;
 - `packmate-prod` contains no Workbench, Playground, or Pipeline.
- REQUIRED EVIDENCE: Argo CD Synced/Healthy, passing `make verify-prod` , and a successful PROD response.

WORKSHOP COMPLETE

You completed the Packmate path from AI prototype to controlled production delivery:

- used a shared OpenShift AI Llama model;
- used Weather and Baggage Policy MCP tools;
- validated AI behavior in the Playground;
- validated the integrated React/FastAPI application in DEV;
- ran an AI-aware Tekton Pipeline with a 0.90 quality gate;
- published an immutable GHCR image reference;
- promoted the validated digest through a reviewed pull request;
- manually synchronized PROD through Argo CD;
- verified the exact production image and application behavior.

You preserved the central platform boundaries: the canonical repository stayed read-only, credentials stayed out of Git, CI produced an immutable candidate, Git recorded the approved desired state, and Argo CD remained the deployment owner.

APPENDIX A — Participant troubleshooting

A.1 Canonical repository promotion is blocked

SYMPTOM: `BLOCKED_CANONICAL_REPOSITORY_PROMOTION` appears, or the fork URL is wrong.

WHAT IT MEANS: `origin` or `GIT_REPO_URL` points to canonical upstream.

RECOVERY: Run `git remote set-url origin https://github.com/YOUR_USERNAME/packmate-agent.git`, then `make configure-participant` and `make verify-demo-fork`.

A.2 Branch settings disagree

SYMPTOM: The current branch, GitOps revision, and promotion base do not match.

WHAT IT MEANS: The workshop branch settings are not coordinated.

RECOVERY: Before baseline preparation, run `git switch packmate-v2` and `make configure-participant`. For the standard prepared path, run `make prepare-demo-baseline`.

A.3 Git push authentication fails

SYMPTOM: A Git push fails, including an `askpass ECONNREFUSED` error.

WHAT IT MEANS: The Workbench cannot reach a valid Git credential prompt, or your credential lacks access to your fork.

RECOVERY: Run `unset GIT_ASKPASS SSH_ASKPASS`, complete `gh auth login` in the terminal, then run `make verify-github-write-readiness`.

A.4 OpenShift is using a workload identity

SYMPTOM: `BLOCKED_OPENSHIFT_SERVICE_ACCOUNT_IDENTITY` appears.

WHAT IT MEANS: `oc` is authenticated as a ServiceAccount instead of your human user.

RECOVERY: Run `oc logout`, then `oc login --web`, and confirm a human username with `oc whoami`.

A.5 The GHCR publication Secret is missing

SYMPTOM: `packmate-ghcr-push` is missing.

WHAT IT MEANS: Tekton cannot mount credentials for candidate publication.

RECOVERY: Run `make configure-promotion-registry` and `make verify-promotion-registry`, then start a new PipelineRun.

A.6 Candidate publication fails

SYMPTOM: Pipeline task `publish-candidate` fails.

WHAT IT MEANS: The GHCR credential is missing, expired, or lacks package-write permission.

RECOVERY: Run `make diagnose-latest-pipelinerun`, renew access with `make configure-promotion-registry`, then start a new PipelineRun.

A.7 There is no promotion difference

SYMPTOM: `BLOCKED_NO_PROMOTION_DIFF` appears.

WHAT IT MEANS: PROD Git already contains the same candidate, so there is no reviewable change.

RECOVERY: Run `make prepare-demo-baseline`, start a new PipelineRun from the prepared branch, then promote that run.

A.8 Argo CD does not show the merged change

SYMPTOM: The application remains Synced after the pull request is merged.

WHAT IT MEANS: The Application has not refreshed the watched fork branch.

RECOVERY: In Argo CD, click **Refresh**. Confirm the repository and revision with `make verify-demo-fork-live`, then wait for OutOfSync.

A.9 PROD verification fails before Sync

SYMPTOM: `make verify-prod` fails before manual Sync.

WHAT IT MEANS: Git contains the approved image, but PROD still runs the previous state.

RECOVERY: Return to Argo CD, manually **Sync**, wait for Synced/Healthy, then run `make verify-prod` again.

A.10 Model or MCP assets are missing

SYMPTOM: The model or one of the MCP assets does not appear ready.

WHAT IT MEANS: Bootstrap did not finish, or the OpenShift AI dashboard cache is stale.

RECOVERY: Run `make bootstrap` and `make verify-dev`, hard-refresh Gen AI studio, and select project `packmate-lab`.

A.11 The Pipeline cannot find test requirements

SYMPTOM: The Pipeline test task cannot find `requirements-dev.txt`.

WHAT IT MEANS: The run used Empty Directory instead of a shared PVC.

RECOVERY: Start a new run and choose **VolumeClaimTemplate, 2 GiB**, for workspace `source`.

A.12 Argo CD access is denied after bootstrap

SYMPTOM: Argo CD denies access after bootstrap completes.

WHAT IT MEANS: Your current SSO token does not include the new group membership.

RECOVERY: Log out of Argo CD, log in via OpenShift again, then reopen `packmate-prod`.

SECURITY RECOVERY RULE If a credential was pasted into a file or Git history, stop. Revoke it in GitHub, remove it from the local file without committing, obtain a replacement, and use only the supported hidden prompt or credential flow.

APPENDIX B — Participant command reference

Run these commands from `/opt/app-root/src/packmate-agent`.

COMMAND	PURPOSE	SAFE TO REPEAT
<code>make configure-participant</code>	Infer fork/branch/GHCR owner and enforce upstream push protection	Yes
<code>make configure-git</code>	Set repository-local commit identity from <code>PACKMATE_GIT_NAME</code> and <code>PACKMATE_GIT_EMAIL</code>	Yes
<code>make verify-demo-fork</code>	Check fork ownership, origin, branch, and canonical protection before bootstrap	Yes, read-only
<code>make verify-github-write-readiness</code>	Check safe GitHub fork authentication and diagnose askpass failures	Yes, read-only by default
<code>make prepare-demo-baseline</code>	Prepare and push the disposable fork branch, then save coordinated config	Yes, fork-only
<code>make configure-promotion-registry</code>	Store GHCR access in OpenShift Secrets through an interactive safe flow	Yes; explicitly updates the Secret
<code>make verify-promotion-registry</code>	Confirm the Pipeline publication Secret and portable baseline	Yes, read-only
<code>make preflight</code>	Check human identity, cluster services, project, model, images, GitOps, and Pipeline prerequisites	Yes, read-only except a temporary probe pod
<code>make bootstrap</code>	Idempotently prepare AI assets, Tekton, Argo CD, DEV, and PROD prerequisites	Yes
<code>make verify-dev</code>	Verify the model, MCP assets, DEV workloads, Routes, health, streaming, and Pipeline	Yes, non-destructive
<code>make verify-gitops</code>	Verify Argo CD ownership, fork source, RBAC, and manual PROD policy	Yes, read-only
<code>make verify-demo-fork-live</code>	Confirm live Argo CD Applications watch your fork and prepared branch	Yes, read-only
<code>make diagnose-latest-pipelinerun</code>	Explain common Pending or Failed Pipeline causes	Yes, read-only
<code>make verify-demo-baseline</code>	Confirm a reviewable promotion difference can exist	Yes, read-only
<code>make promote PIPELINERUN=packmate-ci-example</code>	Validate a successful run and open a one-file digest promotion PR in your fork; replace the example run name first	Creates a fork branch and PR
<code>make verify-prod</code>	Verify PROD after manual Argo CD Sync	Yes, non-destructive

Never use participant commands to push to canonical upstream, store tokens in `config/sandbox.env`, directly apply runtime overlays, update the PROD Deployment image, or synchronize PROD from the Pipeline.